

Every commit, in order

Every commit summary message in this repository, oldest first, grouped by the date it was made. Summary line only — no bodies, no diffs. Appended as new commits land, not written once and frozen.

COMMITTS

138, across 23 days (2026-07-01 through 2026-08-10)

SOURCE

git log --date=format: '%Y-%m-%d' --pretty=format: '%ad|%' --reverse

BRANCH

main

2026-07-01	2026-07-02	2026-07-07	2026-07-10	2026-07-13	2026-07-16
2026-07-17	2026-07-18	2026-07-19	2026-07-20	2026-07-21	2026-07-22
2026-07-23	2026-07-24	2026-07-27	2026-07-29	2026-07-30	2026-07-31
2026-08-04	2026-08-05	2026-08-06	2026-08-07	2026-08-10	

2026-07-01

1 COMMIT

1 Initial commit: Dockersized Postres and SQL Server setup

2026-07-02

2 COMMITTS

1 Schema Test File
2 mssql_local password change

2026-07-07		2 COMMITS
1	Add idempotent seed data for customers and orders	
2	Add SQL Server schema and seed for customers/orders	

2026-07-10		1 COMMIT
1	Trying to reset PASS	

2026-07-13		1 COMMIT
1	Fix credential exposure and seed idempotency bugs found in audit	

2026-07-16		4 COMMITS
1	Ignore __pycache__ in .gitignore	
2	Add schema-discovery tools and register them in the agent (pre-loop-generalization)	
3	Generalize the tool-use loop to handle chained tool calls	
4	Harden agent with read-only DB role, SQL guard, and query logging (Phase 5)	

2026-07-17

4 COMMITS

1

Containerize the agent as a docker-compose service (Phase 6)

2

Fix multi-tool-call bug, cap the tool-use loop, handle API errors

3

Fix diamond text overflow across diagrams, add hardening-pass diagram

4

Reorganize PROJECT_DIAGRAMS into per-build subfolders

2026-07-18

3 COMMITS

1

Add raw ingestion script for allocations CSV (Build 2, Phase 1)

2

Add raw_schema.sql for the raw.allocations staging table

3

Add Build 2, Phase 2 flow diagram (profile & transform)

2026-07-19

3 COMMITS

1

Number the Postgres SQL files to encode their run order

2

Add pandas dependency and the raw-data profiling script

3

Add command-reference cheat sheet as a living project doc

1	Add clean.allocations schema and transformation script (Build 2, Phase 2)
2	Fix cheat sheet's docker exec redirection to PowerShell syntax
3	Add Build 2, Phase 3 flow diagram (agent integration & data quality)
4	Connect the agent to clean.allocations and add a pandera/pytest suite (Build 2, Phase 3)
5	Append Build 2, Phase 3 commands to the living cheat sheet
6	Add Build 2.5, Phase 1 flow diagram (provision the scratch workspace)
7	Provision the agent scratch workspace (Build 2.5, Phase 1)
8	Add save_dataframe, a guarded write tool for agent_scratch (Build 2.5, Phase 2)
9	Wire save_dataframe into the agent and prove a save-then-read-back round trip (Build 2.5, Phase 3)
10	Add automated boundary tests for the agent scratch workspace (Build 2.5, Phase 4)
11	Add standalone Build 2.5 handoff brief
12	Cap the number of distinct tables in agent_scratch (Build 2.5 follow-up)
13	Update handoff brief: table-count cap limitation now resolved

1

Add Build 3, Phase 1 flow diagram (enable semantic search)

2

Rename profile.py to profile_raw_data.py (Build 3 incident)

3

Enable pgvector and backfill embeddings for allocation summaries (Build 3, Phase 1)

4

Fix cheat sheet's profile.py reference after the Build 3 rename

5

Add search_summaries.py, proving raw similarity search works (Build 3, Phase 1)

6

Wire search_summaries into the agent (Build 3, Phase 2)

7

Add Build 3 agent request-flow diagram (runtime lifecycle, not a build phase)

8

Add standalone Build 3 handoff brief (in progress, Phases 1-2)

9

Add automated test suite for semantic search (Build 3, Phase 3)

10

Make embed_summaries.py idempotent - only embeds rows missing summary_embedding (Build 3 follow-up)

11

Update Build 3 handoff brief: re-embedding limitation now resolved

12

Add HNSW vector index on clean.allocations.summary_embedding (Build 3 follow-up)

13

Update Build 3 handoff brief: HNSW vector index limitation now resolved

14

Catch the living cheat sheet up through Build 2.5 and Build 3

15

Add Build 3.5 Phase 1 flow diagram: Chunk & Embed the Docs

16

Add a project-wide file map to the living cheat sheet

1	Add a running list of Claude's own working commands to the cheat sheet
2	Add Build 2 handoff brief — Raw-to-Queryable Data Pipeline
3	Complete Build 3.5 — Documentation RAG Agent (all 3 phases)
4	Add Build 3.5 handoff brief — Documentation RAG Agent
5	Catch the living cheat sheet up through Build 3.5
6	Add missing per-phase flow diagrams for Build 3 and Build 3.5
7	Add Build 4 Phase 1 flow diagram: Structured Tool-Call Tracing
8	Complete Build 4, Phase 1 — structured tool-call tracing
9	Add Build 4 Phase 2 flow diagram: The Eval Harness

1	Build 4 Phase 2: eval harness, plus turn-budget and keyword-robustness fixes
2	Add Build 4 Phase 3 flow diagram: The Metrics Rollup
3	Add question_id to the tool-call log, ahead of Build 5/6 concurrency
4	Document non-obvious design decisions across the project as code comments
5	Build 4 Phase 3: metrics rollup script
6	Add Build 4 handoff brief
7	Catch the living cheat sheet up through Build 4
8	Build 4 follow-up: metrics history, log rotation, and real sample data in the brief
9	Redact leaked dev credentials from Build 1 handoff brief
10	Build 5 Phase 1: CI workflow, masked data fixture, reproducibility fixes

2026-07-24

8 COMMITS

1

Fix CI: scope pytest discovery, remove stray api_test.py

2

Build 5 Phase 3: gated eval harness job in CI

3

Comment-only touch to trigger Phase 3's gated eval job live

4

Replace real+fixture split with one official synthesized dataset

5

Redact real department code from Build 3 handoff brief

6

Skip the genuinely CI-incompatible eval_cases_round2 case explicitly

7

Add Build 5 handoff brief, catch the living cheat sheet up through Build 5

8

Gitignore the personal working-style reference folder

2026-07-27

2 COMMITS

1

Add Build 6 Phase 1 - minimal Streamlit chat UI for the agent

2

Document the Build 6 Phase 1 process incident and next-session handoff prompt

2026-07-29

4 COMMITS

1

Publish Build 6 Phase 2 flow diagram - hybrid tiered conversation memory

2

Cap run_sql_query result size and log the actual question asked

3

Add Build 3.5/4/5 handoff briefs to the docs.chunks source list

4

Build 6 Phase 2 (steps 1-7) - hybrid tiered conversation memory

2026-07-30

5 COMMITS

1

Log the agent's final answer, not just tool-call activity

2

Narrow agent scope to allocations, add one evidence-based prompt fix

3

Publish Build 6 Phase 3 flow diagram - containerize the app

4

Build 6 Phase 3 - containerize the app as a docker-compose service

5

Add a process-ID diagnostic to confirm server restarts

2026-07-31

9 COMMITS

1

Close Build 6 Phase 1 hands-on gate: retype chat UI in place

2

Publish Build 6 Phase 3.5 flow diagram (design deferred, not built)

3

Build 6 Phase 4: surface tool-call activity in a Tools Used expander

4

Phase 4 follow-up: always show the Tools Used expander, even at 0

5

Build 6 handoff brief + living cheat sheet catch-up

6

Build 6 Phase 3.5: persistent shared chat history + FAQ-style expander UI

7

Update all Build 6 docs for Phase 3.5, fix docs.chunks recurring staleness bug

8

Cap agent_scratch.chat_rounds - no display/storage limit existed

9

Document the chat_rounds retention cap fix in both docs, refresh corpus

2026-08-04

4 COMMITS

1

Build 7: verify-before-surface (all 4 phases)

2

Add Build 7 flow diagrams (all 4 phases)

3

Build 7 handoff brief + living cheat sheet catch-up

4

Purge Claude co-author trailer from all history, update cheat sheet convention row

2026-08-05

5 COMMITS

1

Fix CI: add missing 012_chat_rounds_schema.sql to both migration lists

2

Repo audit + full reorganization: src/, migrations/, tests/, docs/, legacy/

3

Fix run_eval.py crashing on log_final_answer/log_verification entries

4

Watch the eval runners themselves in the eval paths-filter

5

Document the two eval-harness bugs the reorg's CI run itself surfaced

2026-08-06

1 COMMIT

1

Track token usage alongside tool use and latency, logged and displayed

1

Fix verifier false positive on subset-condition SQL comparisons

2

Cap individual field size in SQL results, stabilize the eval harness

3

Bake the embedding model into the image, fix cold-start latency

4

Normalize checkout items into `clean.allocation_items`

5

Document the 2026-08-06 token-burn/verifier incident

6

Fix stale `docs.chunks` path and two ambiguous eval cases

7

Fix two more `docs.chunks` retrieval gaps found the same day

8

Add a tracked pre-commit hook to auto-refresh the docs search index

9

Fix `.githooks/pre-commit` missing its executable bit in git's index

10

Catch the cheat sheet up through the 2026-08-07 incident

11

Add the 2026-08-07 deployment-verification incident audit

12

Cross-reference the 2026-08-07 incident, fix a stale claim

13

Fix an exponential history-compounding bug and a docs-routing gap

14

Fold two more same-day findings into the 2026-08-07 incident audit

15

Cap the fidelity window's payload size, fix an AND/OR bug, stop two silent failures

16

Fold Phase 9 and Phase 10 into the 2026-08-07 incident audit

17

Generalize the grounding rule from `search_docs` to every evidence tool

18

Close Phase 11 in the 2026-08-07 audit, evidenced with a live screenshot

1

Route open/deferred-item questions to data, not search, and stop chasing truncated evidence

2

Add the 2026-08-10 stale-status-claim incident audit

3

Split long chunks at extraction time - the mechanism, not the instance

4

Add Phase 5 to the 2026-08-10 audit: the logic fix, not just the situation

5

Close the scratch-table cleanup deferral with age-based tracking

6

Close out the Known Open & Deferred Items table

7

Add immediate scratch-table deletion, bypassing the age gate

8

Document immediate scratch-table deletion in the cheat sheet

9

Fix a subset-vs-total miscount, add the commit log as a new living doc

10

Catch the commit log up through its own addition commit

11

Show per-tool-call result size next to latency, to isolate token burn

12

Catch the commit log up through the tool-call-size feature

13

Rewrite the project overview, swap it into the RAG corpus, and teach the verifier to distinguish contradiction from unconfirmable

2026-08-11

4 COMMITS

1

Link the README to the new portfolio site

2

Add a public-demo entry point with session-scoped history and a hard daily token cap

3

Link the live public demo from the README

4

Replace the verifier's boolean+prose-prefix convention with a verdict enum

Living document — one row per commit, summary line only, grouped by calendar date. Regenerated from `git log` as new commits land, not written once and frozen. See [sql-test-01-cheatsheet.html](https://github.com/ethereum/eth2-testnet/blob/master/testnet-01/cheatsheet.html) for what each change actually did.